



Lecture 7: MDPs I



Final project: task definition

- **Topic:** anything you're excited about (games, music, recommendation systems, optimizing society, personal productivity); see website for ideas; must be well-motivated
- **Scope:** not too broad (automatically replying to emails), not too narrow (training logistic regression for spam classification)
- **Evaluation metric:** quantitative measure of success; speed? accuracy? user studies?
- **Dataset:** start with 3-10 manual examples; then either find existing dataset or scrape new dataset; make sure you have data!

Final project: approach

- Baseline: simple method (logistic regression with SGD); reveals challenges to tackle; provides lower bound on performance
- Oracle: cheating method that is unrealistically optimistic (if trying to speed up X, assume X takes zero time); provides upper bound on performance
- Main approach: develop approaches that make sense for the task; try to keep things simple

Final project: analysis

- Think like a scientist: ask questions and answer them (e.g., does your method work better than the baseline when there are few training examples?)
- Error analysis: show concrete examples and categorize errors, suggest future work
- Literature review: compare with other pieces of work at a technical level; simulate them if possible (e.g., if they don't use context, then strip out context from your approach)

Final project: example task definition

- **Task:** automatically illustrate a news article with images (motivation: pictures make reading more fun)
- **Scope:** focus on single paragraphs from Google News
- **Evaluation metric:** given two paragraphs and their two correct images, figure out which one is which
- **Dataset:** scrape 1000 articles from Google News

Final project: example approach

- **Baseline:** for each image, run standard object detector; return image with highest overlap with words in paragraph
- **Oracle:** have humans do it; or hand-label paragraph with relevant words and desired objects
- **Algorithms:** (i) predict an object (or none) for each word in the paragraph, perform word similarity with objects detected from images; or (ii) Google image search for related images of a word and perform object similarity

Final project: example analysis

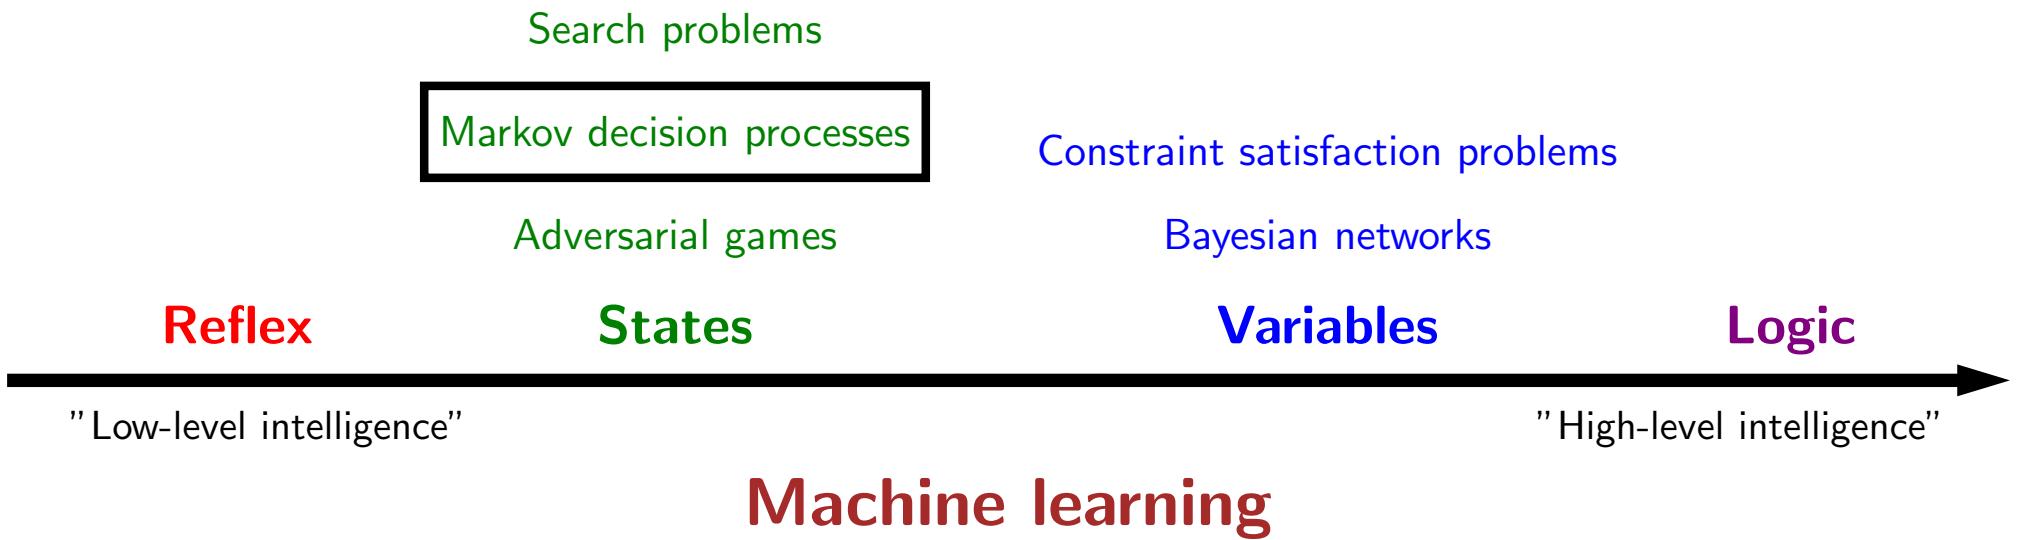
- **Question:** which is more reliable, word similarity or object similarity? word similarity is easier to work with
- **Question:** which types of articles are easier? concrete events easier than abstract topics
- **Question:** was the model able to figure out which words are relevant? common nouns are easy, but sometimes spurious correlations get in the way

Final project

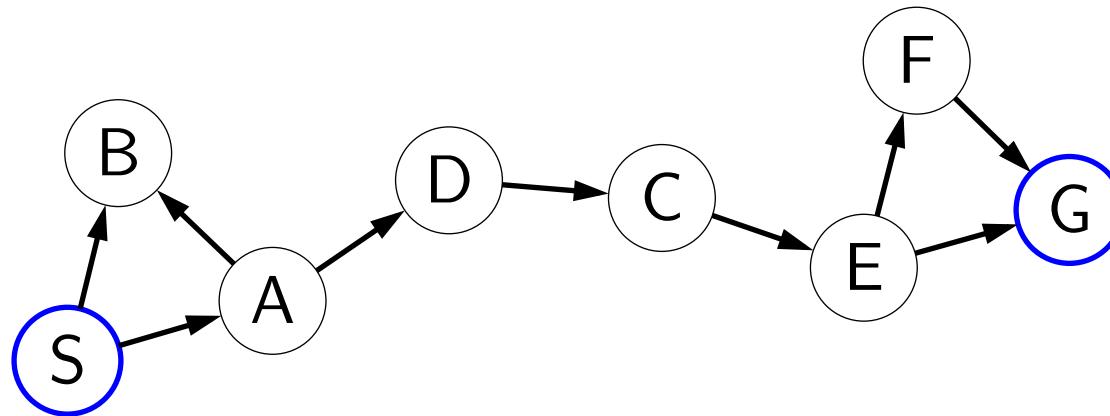
- **Summary:** task definition, approach, analysis
- **Project proposal:** due next Tuesday (Oct. 21)
- **Get help:** come to office hours!

Questions?

Course plan



So far: search problems

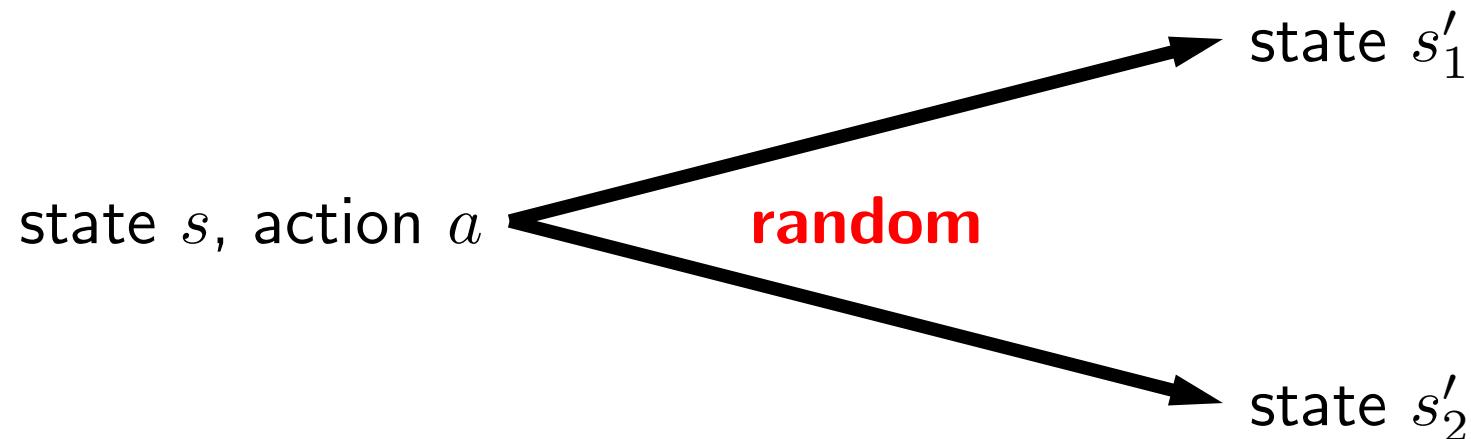


state s , action a $\xrightarrow{\text{deterministic}}$ state $\text{Succ}(s, a)$



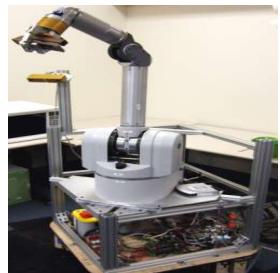
- Last week, we looked at search problems, a powerful paradigm that can be used to solve a diverse range of problems ranging from word segmentation to package delivery to route finding. The key was to cast whatever problem we were interested in solving into the problem of finding the minimum cost path in a graph.
- However, search problems assume that taking an action a from a state s results **deterministically** in a unique successor state.

Uncertainty in the real world



- In the real world, the deterministic successor assumption is often unrealistic, for there is **randomness**: taking an action might lead to any one of many possible states.
- Today, we will develop tools to tackle this more challenging setting. We will fortunately still be able to reuse many of the intuitions about search problems, in particular the notion of a state.

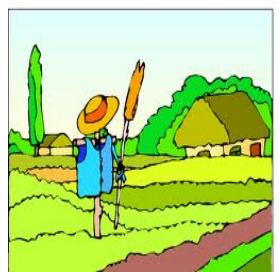
Applications



Robotics: decide where to move, but actuators can fail, hit unseen obstacles, etc.



Resource allocation: decide what to produce, don't know the customer demand for various products



Agriculture: decide what to plant, but don't know weather and thus crop yield

- Randomness shows up in many places. They could be caused by limitations of the sensors and actuators of the robot (which we can control to some extent). Or they could be caused by market forces or nature, which we have no control over.
- We'll see that all of these sources of randomness can be handled in the same mathematical framework.

Volcano crossing



Run (or press ctrl-enter)

		-50	20
		-50	
2			

- Let us consider an example. You are exploring a South Pacific island, which is modeled as a 3x4 grid of states. From each state, you can take one of four actions to move to an adjacent state: north (N), east (E), south (S), or west (W). If you try to move off the grid, you remain in the same state. You start at (2,1). If you end up in either of the green or red squares, your journey ends, either in a lava lake (reward of -50) or in a safe area with either no view (2) or a fabulous view of the island (20). What do you do?
- If we have a deterministic search problem, then the obvious thing will be to go for the fabulous view, which yields a reward of 20. You can set `numIters` to 10 and press Run. Each state is labeled with the maximum expected utility (sum of rewards) one can get from that state (analogue of FutureCost in a search problem). We will define this quantity formally later. For now, look at the arrows, which represent the best action to take from each cell. Note that in some cases, there is a tie for the best, where some of the actions seem to be moving in the wrong direction. This is because there is no penalty for moving around indefinitely. If you change `moveReward` to -0.1, then you'll see the arrows point in the right direction.
- In reality, we are dealing with treacherous terrain, and there is on each action a probability `slipProb` of slipping, which results in moving in a random direction. Try setting `slipProb` to various values. For small values (e.g., 0.1), the optimal action is to still go for the fabulous view. For large values (e.g., 0.3), then it's better to go for the safe and boring 2. Play around with the other reward values to get intuition for the problem.
- Important: note that we are only specifying the dynamics of the world, not directly specifying the best action to take. The best actions are computed automatically from the algorithms we'll see shortly.



Roadmap

Markov decision process

Policy evaluation

Policy iteration

Value iteration

Dice game



Example: dice game

For each round $r = 1, 2, \dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
 - If the dice results in 1 or 2, we end the game.
 - Otherwise, continue to the next round.

Start

Stay

Quit

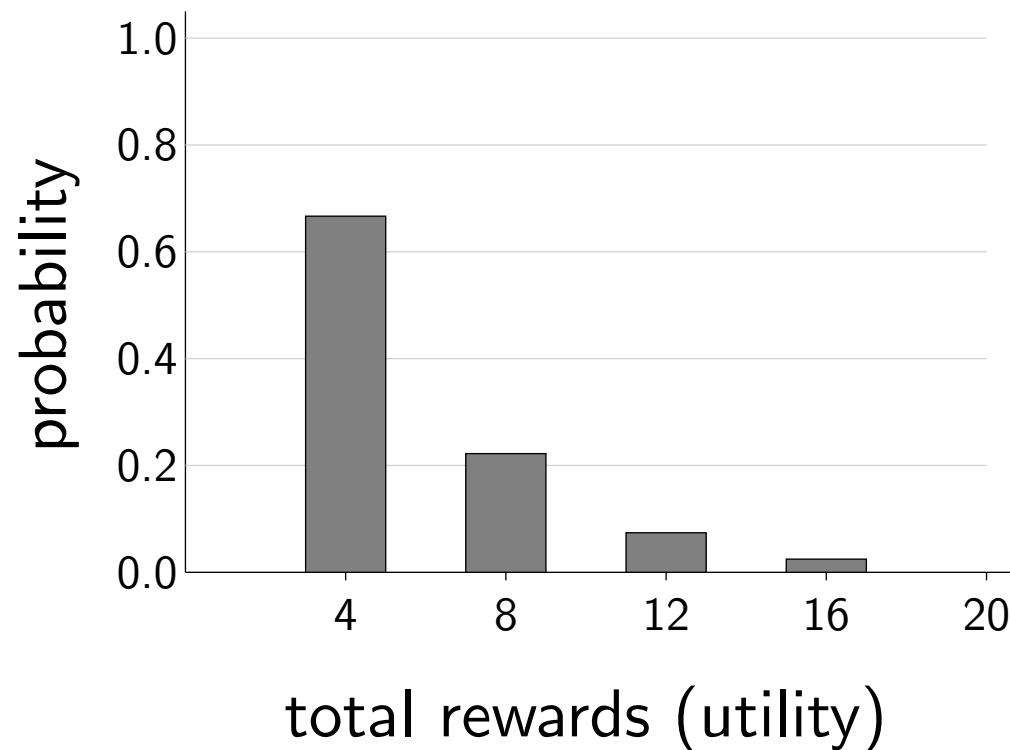
Dice:

Rewards:
 0

- We'll see more volcanoes later, but let's start with a much simpler example: a dice game. What is the best strategy for this game?

Rewards

If follow policy "stay":



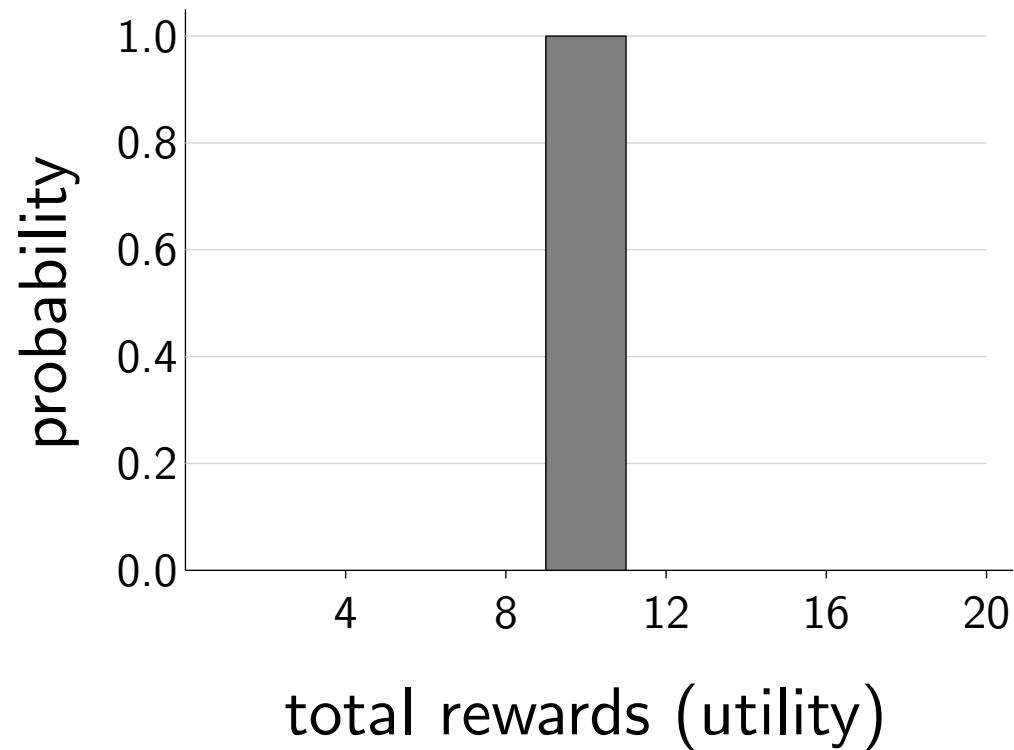
Expected utility:

$$\frac{1}{3}(4) + \frac{2}{3} \cdot \frac{1}{3}(8) + \frac{2}{3} \cdot \frac{2}{3} \cdot \frac{1}{3}(12) + \dots = 12$$

- Let's suppose you always stay. Note that each outcome of the game will result in a different sequence of rewards, resulting in a **utility**, which is in this case just the sum of the the rewards.
- We are interested in the **expected** utility, which you can compute to be 12.

Rewards

If follow policy "quit":



Expected utility:

$$1(10) = 10$$

- If you quit, then you'll get a reward of 10 deterministically. Therefore, in expectation, the "stay" strategy is preferred, even though sometimes you'll get less than 10.

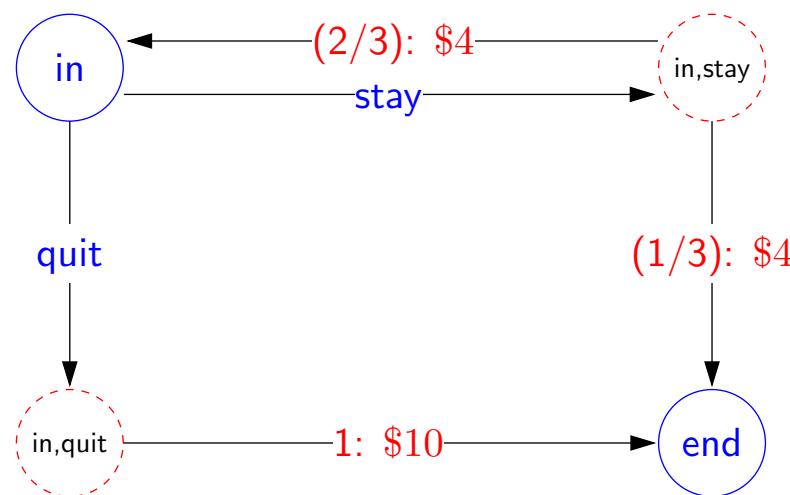
MDP for dice game



Example: dice game

For each round $r = 1, 2, \dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
 - If the dice results in 1 or 2, we end the game.
 - Otherwise, continue to the next round.



- While we already solved this game directly, we'd like to develop a more general framework for thinking about not just this game, but also other problems such as the volcano crossing example. To that end, let us formalize the dice game as a **Markov decision process** (MDP).
- An MDP can be represented as a graph. The nodes in this graph include both **states** and **chance nodes**. Edges coming out of states are the possible actions from that state, which lead to chance nodes. Edges coming out of a chance nodes are the possible random outcomes of that action, which end up back in states. Our convention is to label these chance-to-state edges with the probability of a particular **transition** and the associated reward for traversing that edge.

Markov decision process



Definition: Markov decision process

States: the set of states

$s_{\text{start}} \in \text{States}$: starting state

$\text{Actions}(s)$: possible actions from state s

$T(s, a, s')$: probability of s' if take action a in state s

$\text{Reward}(s, a, s')$: reward for the transition (s, a, s')

$\text{IsEnd}(s)$: whether at end of game

$0 \leq \gamma \leq 1$: discount factor (default: 1)

- A **Markov decision process** has a set of states States , a starting state s_{start} , and the set of actions $\text{Actions}(s)$ from each state s .
- It also has a **transition distribution** T , which specifies for each state s and action a , a distribution over successor states s' .
- Associated with each transition (s, a, s') is a reward, which could be either positive or negative.
- If we arrive in a state s for which $\text{IsEnd}(s)$ is true, then the game is over.
- Finally, the discount factor γ is a quantity which specifies how much we value the future and will be discussed later.

Search problems



Definition: search problem

States: the set of states

$s_{\text{start}} \in \text{States}$: starting state

$\text{Actions}(s)$: possible actions from state s

$\text{Succ}(s, a)$: where we end up if take action a in state s

$\text{Cost}(s, a)$: cost for taking action a in state s

$\text{IsGoal}(s)$: whether at end

- $\text{Succ}(s, a) \Rightarrow T(s, a, s')$
- $\text{Cost}(s, a) \Rightarrow \text{Reward}(s, a, s')$
- $\text{IsGoal}(s) \Rightarrow \text{IsEnd}(s)$

- MDPs share many similarities with search problems, but there are differences (one main difference and two minor ones).
- The main difference, is the move from a deterministic successor function $\text{Succ}(s, a)$ to transition probabilities over s' . We can think of the successor function $\text{Succ}(s, a)$ as a special case of transition probabilities:

$$T(s, a, s') = \begin{cases} 1 & \text{if } s' = \text{Succ}(s, a) \\ 0 & \text{otherwise} \end{cases}.$$

- A minor difference is that we've gone from minimizing costs to maximizing rewards. The two are really equivalent: you can negate one to get the other.
- A second minor difference is that search problems have a goal test $\text{IsGoal}(s)$, but MDPs have a function $\text{IsEnd}(s)$ that tests whether we are at the end. The use of $\text{IsEnd}(s)$ is just to deemphasize the positive connotation associated with reaching a goal; there could be both bad and good endings.

Transitions



Definition: transition probabilities

The **transition probabilities** $T(s, a, s')$ specify the probability of ending up in state s' if taken action a in state s .



Example: transition probabilities

s	a	s'	$T(s, a, s')$
in	quit	end	1
in	stay	in	2/3
in	stay	end	1/3

- Just to dwell on the major difference, transition probabilities, a bit more: for each state s and action a , the transition probabilities specifies a distribution over successor states s' .

Probabilities sum to one



Example: transition probabilities

s	a	s'	$T(s, a, s')$
in	quit	end	1
in	stay	in	2/3
in	stay	end	1/3

For each state s and action a :

$$\sum_{s' \in \text{States}} T(s, a, s') = 1$$

Successors: s' such that $T(s, a, s') > 0$

- This means that for each given s and a , if we sum the transition probability $T(s, a, s')$ over all possible successor states s' , we get 1.
- If a transition to a particular s' is not possible, then $T(s, a, s') = 0$. We refer to the s' for which $T(s, a, s') > 0$ as the successors.
- Generally, the number of successors of a given (s, a) is much smaller than the total number of states. For instance, in a search problem, each (s, a) has exactly one successor.

What is a solution?

Search problem: path (sequence of actions)

MDP:



Definition: policy

A **policy** π is a mapping from each state $s \in \text{States}$ to an action $a \in \text{Actions}(s)$.



Example: volcano crossing

s	$\pi(s)$
(1,1)	S
(2,1)	E
(3,1)	N
...	...

- So we now know what an MDP is. What do we do with one? For search problems, we were trying to find the minimum cost **path**
- However, fixed paths won't suffice for MDPs, because we don't know which states the random dice rolls are going to take us.
- Therefore, we define a **policy**, which specifies an action for every single state, not just the states along a path. This way, we have all our bases covered, and know what action to take no matter where we are.



Roadmap

Markov decision process

Policy evaluation

Policy iteration

Value iteration

Evaluating a policy



Definition: utility

Following the policy yields a **random path**.

The **utility** of a policy is the (discounted) sum of the rewards on the path (this is a random quantity).

Path	Utility
[in, stay, 4, end]	4
[in, stay, 4, in, stay, 4, in, stay, 4, end]	12
[in, stay, 4, in, stay, 4, end]	8
[in, stay, 4, in, stay, 4, in, stay, 4, in, stay, 4, end]	16
...	...



Definition: value (expected utility)

The **value** of a policy is the **expected utility**.

- Now that we've defined an MDP (the input) and a policy (the output), let's turn to defining the evaluation metric for a policy — there are many of them, which one should we choose?
- Recall that we'd like to maximize the total rewards (utility), but this is a random quantity, so we can't quite do that. Instead, we will instead maximize the **expected utility**, which we will refer to as **value** (of a policy).

Evaluating a policy: volcano crossing

Run

(or press ctrl-enter)

a	r	s	
	(2,1)		
E	-0.1	(2,2)	
S	-0.1	(3,2)	
E	-0.1	(3,3)	
E	-50.1	(2,3)	

2.4	-0.5	-50	40
3.7 →	5	-50	31
2	12.6 →	16.3 →	26.2

Value: 3.73

Utility: -36.79

- To get an intuitive feel for the relationship between a value and utility, consider the volcano example. If you press Run multiple times, you will get random paths shown on the right leading to different utilities. Note that there is considerable variation in what happens.
- The expectation of this utility is the **value**.
- You can run multiple simulations by increasing numEpisodes. If you set numEpisodes to 1000, then you'll see the average utility converging to the value.

Discounting



Definition: utility

Path: $s_0, a_1, r_1, s_1, a_2, r_2, s_2, \dots$ (state, action, reward).

The **utility** with discount γ is

$$u_1 = r_1 + \gamma r_2 + \gamma^2 r_3 + \gamma^3 r_4 + \dots$$

Discount $\gamma = 1$ (save for the future):

$$[\text{stay, stay, stay, stay}]: 4 + 4 + 4 + 4 = 16$$

Discount $\gamma = 0$ (live in the moment):

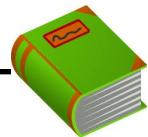
$$[\text{stay, stay, stay, stay}]: 4 + 0 \cdot (4 + \dots) = 4$$

Discount $\gamma = 0.5$ (compromise):

$$[\text{stay, stay, stay, stay}]: 4 + \frac{1}{2} \cdot 4 + \frac{1}{4} \cdot 4 + \frac{1}{8} \cdot 4 = 7.5$$

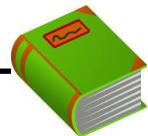
- There is an additional aspect to utility: discounting, which captures the fact that a reward today is worth more than the same reward tomorrow. If the discount γ is small, that means we favor the present more and downweight future rewards more.
- In what's to come, we will not work with the random variables (utilities), but instead work directly with their expectations (values / expected utilities).

Policy evaluation



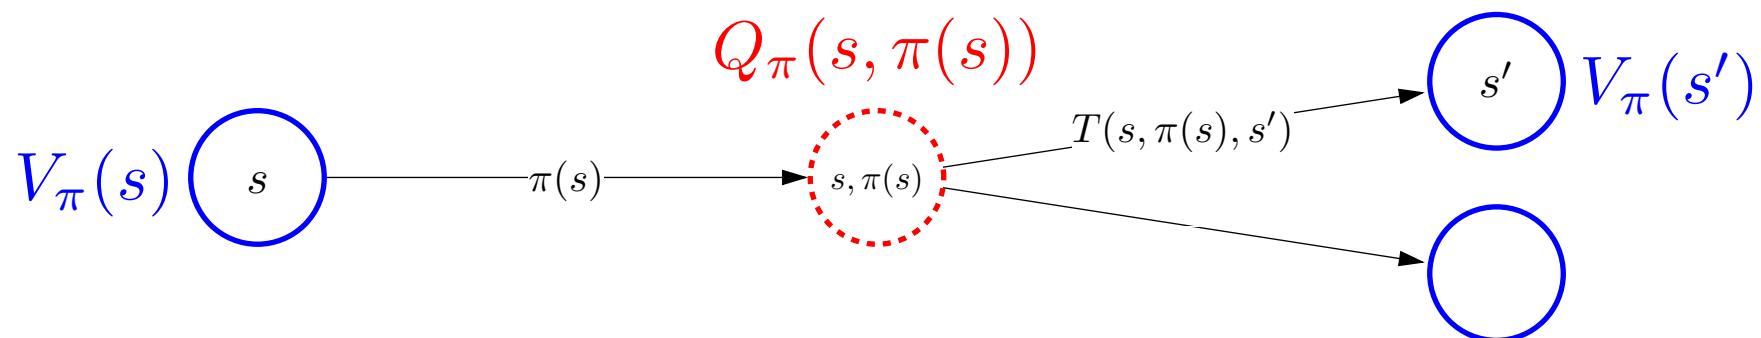
Definition: value of a policy

Let $V_\pi(s)$ be the expected utility received by following policy π from state s .



Definition: Q-value of a policy

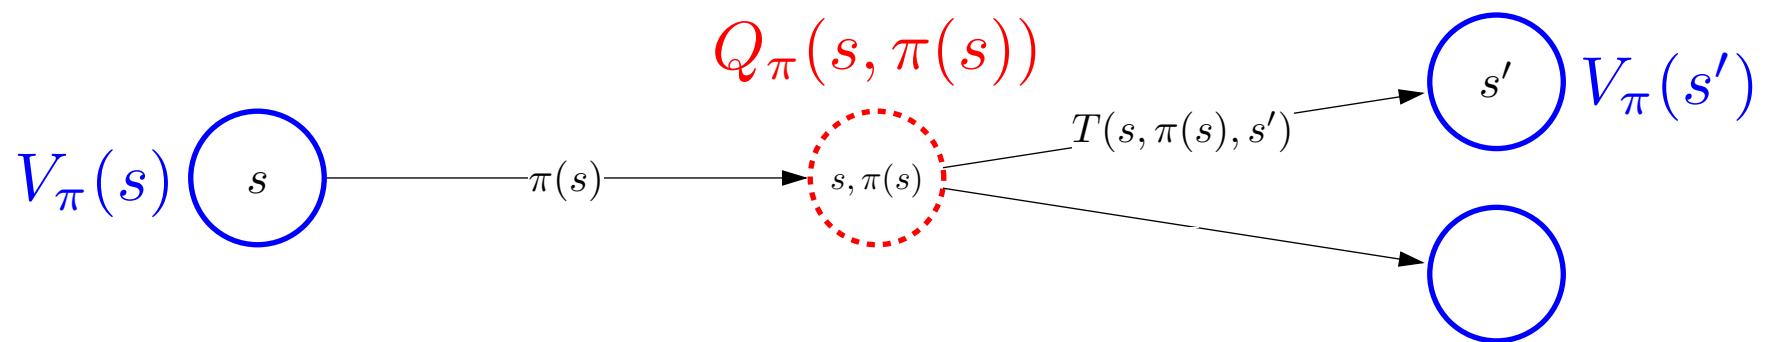
Let $Q_\pi(s, a)$ be the expected utility of taking action a from state s , and then following policy π .



- Associated with any policy π are two important quantities, the value of the policy $V_\pi(s)$ and the Q-value of a policy $Q_\pi(s, a)$.
- In terms of the MDP graph, one can think of the value $V_\pi(s)$ as labeling the state nodes, and the Q-value $Q_\pi(s, a)$ as labeling the chance nodes.
- This label refers to the expected utility if we were to start at that node and continue the dynamics of the game.

Policy evaluation

Plan: define recurrences relating value and Q-value

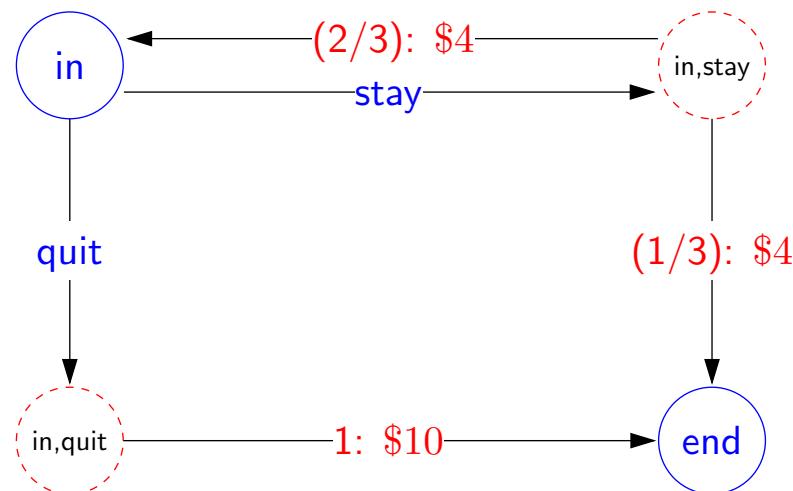


$$V_\pi(s) = \begin{cases} 0 & \text{if } \text{IsEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise.} \end{cases}$$

$$Q_\pi(s, a) = \sum_{s'} T(s, a, s')[\text{Reward}(s, a, s') + \gamma V_\pi(s')]$$

- We will now write down some equations relating value and Q-value. Our eventual goal is to get to an algorithm for computing these values, but as we will see, writing down the relationships gets us most of the way there, just as writing down the recurrence for FutureCost directly lead to a dynamic programming algorithm for acyclic search problems.
- First, $V_\pi(s)$, the value of a state s is gotten by just following the action edge specified by the policy and taking the Q-value $Q_\pi(s, \pi(s))$. (There's also a base case where $\text{IsEnd}(s)$.)
- Second, $Q_\pi(s, a)$ is gotten by considering all possible transitions to successor states s' and taking the expectation over the immediate reward $\text{Reward}(s, a, s')$ plus the discounted future reward $\gamma V_\pi(s')$.

Dice game



Let π be the "stay" policy: $\pi(\text{in}) = \text{stay}$.

$$V_\pi(\text{in}) = Q_\pi(\text{in}, \text{stay}) = \frac{1}{3}(4) + \frac{2}{3}(4 + 1 \cdot V_\pi(\text{in}))$$

Compute in closed form: $V_\pi(\text{in}) = 12$

- As an example, let's compute the values of the nodes in the dice game for the policy "stay".
- Note that the recurrence involves both $V_\pi(\text{in})$ on the left-hand side and the right-hand side. We can solve this recurrence easily to get the value.

Policy evaluation



Key idea: iterative algorithm

Start with arbitrary policy values and repeatedly apply recurrences to converge to true values.



Algorithm: policy evaluation

Initialize $V_{\pi}^{(0)}(s) \leftarrow 0$ for all states s .

For iteration $t = 1, \dots, t_{PE}$:

For each state s :

$$V_{\pi}^{(t)}(s) \leftarrow \underbrace{\sum_{s'} T(s, \pi(s), s')[\text{Reward}(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]}_{Q^{(t-1)}(s, \pi(s))}$$

- We have so far mathematically characterized the value of a policy via the two recurrences. However, how do we actually calculate it?
- One can observe that it's a system of linear equations in $V_\pi(s)$, so any linear equation solver would work. However, we can actually use an even simpler approach called **policy evaluation**.
- We've already seen examples of iterative algorithms in machine learning: the basic idea is to start with something crude, and refine it over time.
- Policy iteration starts with a vector of all zeros for the initial values $V_\pi^{(0)}$. Each iteration, we loop over all the states and apply the two recurrences that we had before. The equations look hairier because of the superscript (t) , which simply denotes the value of at iteration t of the algorithm.

Policy evaluation implementation

How many iterations (t_{PE})? Repeat until values don't change much:

$$\max_{s \in \text{States}} |V_\pi^{(t)}(s) - V_\pi^{(t-1)}(s)| \leq \epsilon$$

Don't store $V_\pi^{(t)}$ for each iteration t , need only last two:

$$V_\pi^{(t)} \text{ and } V_\pi^{(t-1)}$$

- Some implementation notes: a good strategy for determining how many iterations to run policy evaluation is based on how accurate the result is. Rather than set some fixed number of iterations (e.g, 100), we instead set an error tolerance (e.g., $\epsilon = 0.01$), and iterate until the maximum change between values of any state s from one iteration (t) to the previous ($t - 1$) is at most ϵ .
- The second note is that while the algorithm is stated as computing $V_{\pi}^{(t)}$ for each iteration t , we actually only need to keep track of the last two values. This is important for saving memory.

Complexity



Algorithm: policy evaluation

Initialize $V_{\pi}^{(0)}(s) \leftarrow 0$ for all states s .

For iteration $t = 1, \dots, t_{\text{PE}}$:

For each state s :

$$V_{\pi}^{(t)}(s) \leftarrow \underbrace{\sum_{s'} T(s, \pi(s), s')[\text{Reward}(s, \pi(s), s') + \gamma V_{\pi}^{(t-1)}(s')]}_{Q^{(t-1)}(s, \pi(s))}$$

MDP complexity

S states

A actions per state

S' successors (number of s' with $T(s, a, s') > 0$)

Time: $O(t_{\text{PE}} S S')$

- Computing the running time of policy evaluation is straightforward: for each of the t_{PE} iterations, we need to enumerate through each of the S states, and for each one of those, loop over the successors S' . Note that we don't have a dependence on the number of actions A because we have a fixed policy $\pi(s)$ and we only need to look at the action specified by the policy.
- Advanced: Here, we have to iterate t_{PE} time steps to reach a target level of error ϵ . It turns out that t_{PE} doesn't actually have to be very large for very small errors. Specifically, the error decreases exponentially fast as we increase the number of iterations. In other words, to cut the error in half, we only have to run a constant number of more iterations.
- Advanced: For acyclic graphs (for example, the MDP for Blackjack), we just need to do one iteration (not t_{PE}) provided that we process the nodes in reverse topological order of the graph. This is the same setup as we had for dynamic programming in search problems, only the equations are different.

Policy evaluation on dice game

Let π be the "stay" policy: $\pi(\text{in}) = \text{stay}$.

$$V_{\pi}^{(t)}(\text{in}) = \frac{1}{3}(4) + \frac{2}{3}(4 + V_{\pi}^{(t-1)}(\text{in}))$$

s	end	in	$(t = 100 \text{ iterations})$
$V_{\pi}^{(t)}$	0.00	12.00	

Converges to $V_{\pi}(\text{in}) = 12$.

- Let us run policy evaluation on the dice game. The value converges very quickly to the correct answer.



Summary so far

- MDP: graph with states, chance nodes, transition probabilities, rewards
- Policy: mapping from state to action (solution to MDP)
- Value of policy: expected utility over random paths
- Policy evaluation: iterative algorithm to compute value of policy

- Let's summarize: we have defined an MDP, which we should think of a graph where the nodes are states and chance nodes. Because of randomness, solving an MDP means generating policies, not just paths. A policy is evaluated based on its value: the expected utility obtained over random paths. Finally, we saw that policy evaluation provides a simple way to compute the value of a policy.



Roadmap

Markov decision process

Policy evaluation

Policy iteration

Value iteration

Policy improvement

So far: policy evaluation computes value of a fixed policy π

Goal: improve π to something slightly better π_{new}

Recall: $Q_\pi(s, a)$ is the expected utility of first taking action a in state s , and then following π

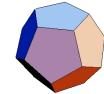


Algorithm: policy improvement

1. Compute $Q_\pi(s, a)$ from $V_\pi(s)$
2. Compute $\pi_{\text{new}}(s) = \arg \max_{a \in \text{Actions}(s)} Q_\pi(s, a)$

- If we are given a policy π , we now know how to compute its value $V_\pi(s_{\text{start}})$. So now, we could just enumerate all the policies, compute the value of each one, and take the best policy, but the number of policies is exponential in the number of states (A^S to be exact), so we need something a bit more clever.
- Let's set a modest goal of just trying to **improve** the current policy π into something a bit better π_{new} . How do we do this? We've done all this work to compute V_π and Q_π ; surely those contain more information that can help us.
- Recall that $Q_\pi(s, a)$ is the value if we took action a from state s , and then followed π . So instead of taking action $\pi(s)$ in state s , why not take the best action according to $Q_\pi(s, a)$: $\arg \max_{a \in \text{Actions}(s)} Q_\pi(s, a)$. This is exactly what policy improvement does.

Policy improvement



Example: dice game

Suppose $\pi(\text{in}) = \text{quit}$.

Step 1:

$$Q_\pi(\text{in}, \text{quit}) = 10.$$

$$Q_\pi(\text{in}, \text{stay}) = 4 + \frac{2}{3}(10) \approx 10.67.$$

Step 2:

$$\pi_{\text{new}}(\text{in}) = \text{stay}$$

- In the dice game, suppose our initial policy is "quit".
- Policy evaluation produces $V_\pi(s)$ for each state.
- In the first step of policy improvement, we compute $Q_\pi(s, a)$ for all states and actions. (Recall that policy evaluation only computed $Q_\pi(s, a)$ for a single action $a = \pi(s)$.)
- In the second step, we choose the best action for every state. In this case, we recognize that the "stay" action yields a higher Q-value than the "quit" action.

Policy improvement



Algorithm: policy improvement

1. Compute $Q_\pi(s, a)$ from $V_\pi(s)$
2. Compute $\pi_{\text{new}}(s) = \arg \max_{a \in \text{Actions}(s)} Q_\pi(s, a)$

-MDP complexity

S states

A actions per state

S' successors (number of s' with $T(s, a, s') > 0$)

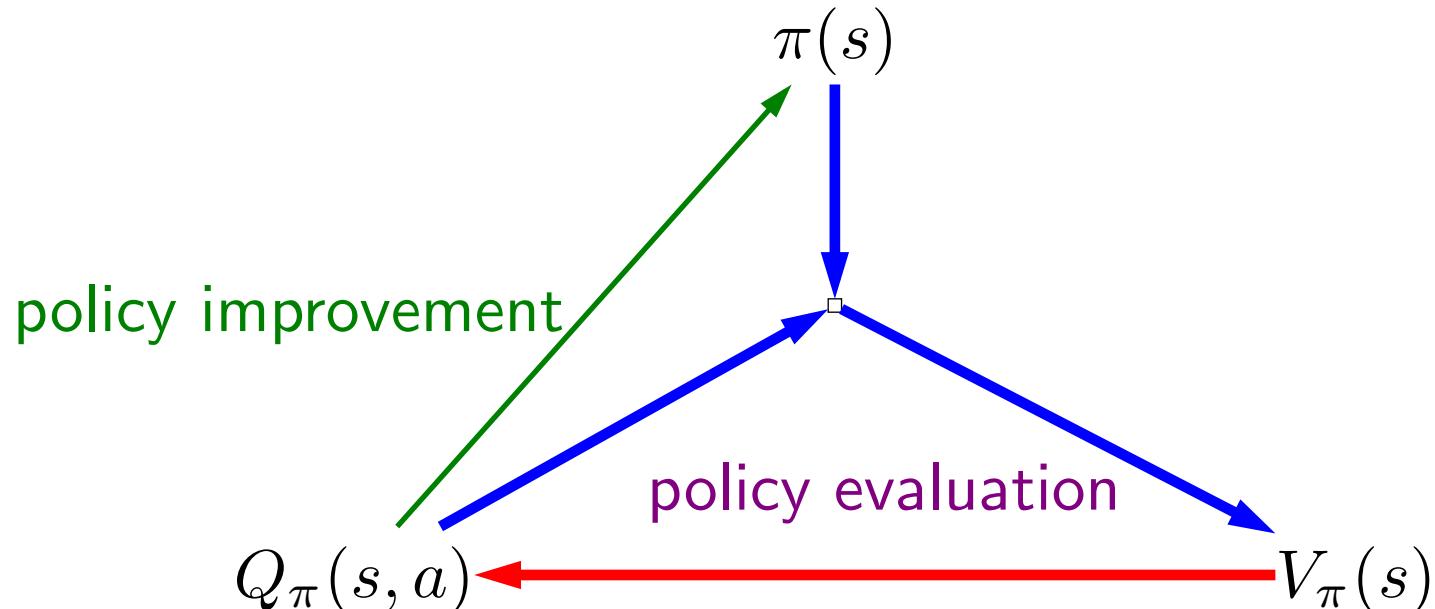
Time: $O(SAS')$

Note: unlike policy evaluation, no iterating to some error tolerance

- Policy improvement is relatively cheap. Unlike policy evaluation, we don't have to iterate to some error tolerance.
- The main bottleneck is actually computing $Q_\pi(s, a)$ in step 1, which takes $O(SAS')$ time. Computing the actual new policy π_{new} only takes $O(SA)$ time.



Connections



Methods

Policy evaluation $(\rightarrow \rightarrow)^*$

Policy improvement \rightarrow

Policy iteration $((\rightarrow \rightarrow)^* \rightarrow)^*$

Value iteration $(\rightarrow \rightarrow \rightarrow)^*$

- So far, we have talked about policy evaluation, which can only tell us how well we're doing if we have a policy, and policy improvement, which allows us to improve the policy.
- Policy evaluation and policy improvement lay the conceptual groundwork for MDP algorithms. The other two more algorithms, policy iteration and value iteration, merely combine the existing ideas.
- We can visualize all the algorithms as simply repeated application of our standard recurrences.
- Policy evaluation iterates between the blue arrow (updating the value of states) and the red arrows (updating Q-values of chance nodes). The * means that we repeatedly apply the updates.
- Policy improvement applies the green arrow once to get a new value of $\pi(s)$.
- This figure is a bit of a simplification. Policy evaluation technically only computes $Q_\pi(s, \pi(s))$, not $Q_\pi(s, a)$. Policy improvement technically starts with $V_\pi(s)$ and computes $Q_\pi(s, a)$ before updating $\pi(s)$. But this is still a useful conceptual diagram to have in one's head.

Policy iteration

Idea: rinse and repeat



Algorithm: policy iteration

$\pi \leftarrow \text{arbitrary}$

For $t = 1, \dots, t_{\text{PI}}$ (or until π stops changing):

 Run policy evaluation to compute V_π

 Run policy improvement to get π_{new}

$\pi \leftarrow \pi_{\text{new}}$

Time: $O(t_{\text{PI}}(t_{\text{PE}}SS' + SAS'))$

Implementation trick: **warm start** policy evaluation with previous V_π

- Policy evaluation allows us to go from a policy π to a value function V_π . Policy improvement allows us to take that value function V_π and generate a new policy π_{new} .
- The natural thing to do now is to repeatedly compute the value of a policy and improve it. This is exactly policy iteration.
- Policy iteration is guaranteed to converge to a policy that obtains the globally optimal value.
- Note that there are two nested loops now: an outer loop for policy iteration and an inner loop for policy evaluation.
- The running time depends on the product of the number of inner iterations and number of outer iterations. The number of inner iterations depends on how close the initialization is to the final result (a general property of iterative algorithms). One common trick, called **warm start**, is to initialize the inner loop with our best guess. If the policy is not changing very much, then the inner policy evaluation loop will converge quickly.



Roadmap

Markov decision process

Policy evaluation

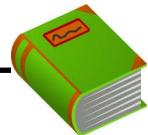
Policy iteration

Value iteration

- So now, we already have a pretty good algorithm, policy iteration, for computing the optimal policy given an MDP. We will now introduce value iteration, which is an alternative algorithm, which gets at the problem more directly. It is closely related to policy iteration.

Optimal value and policy

Goal: try to get directly at maximum expected utility

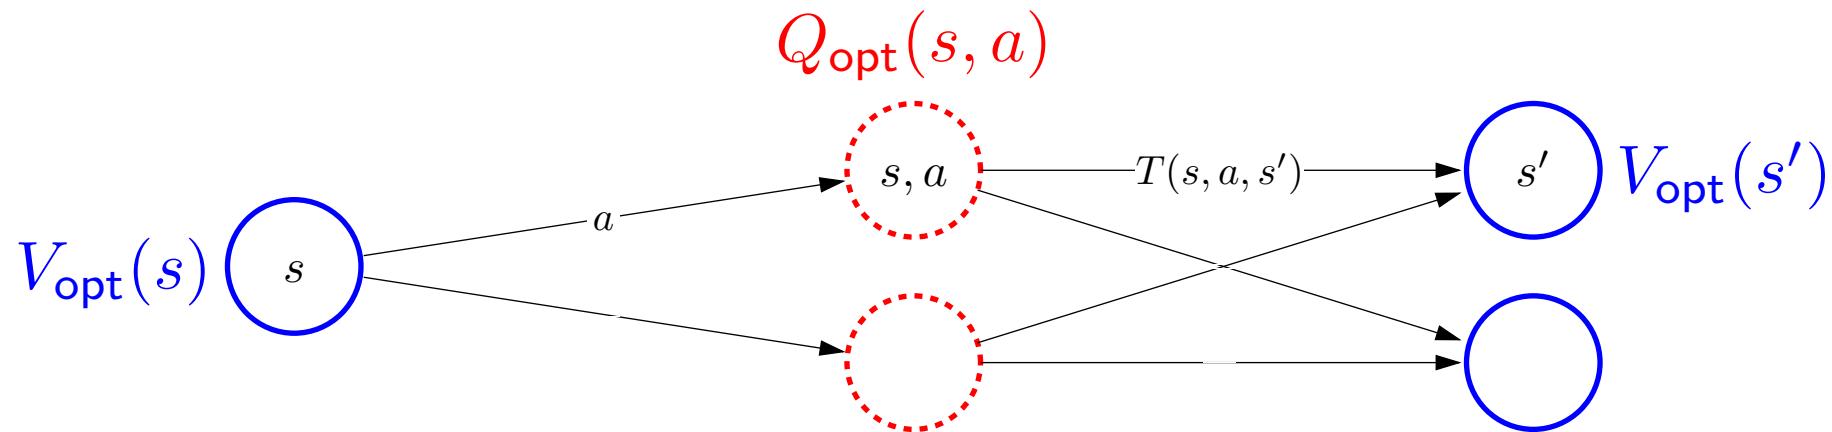


Definition: optimal value

The **optimal value** $V_{\text{opt}}(s)$ is the maximum value attained by any policy.

- The idea is not to worry about having a policy π , but just directly compute the value of the optimal policy. We will write down a bunch of recurrences which look exactly like policy evaluation, but instead of having V_π and Q_π with respect to a fixed policy π , we will have V_{opt} and Q_{opt} , which are with respect to the optimal policy.

Optimal values and Q-values



Optimal value if take action a in state s :

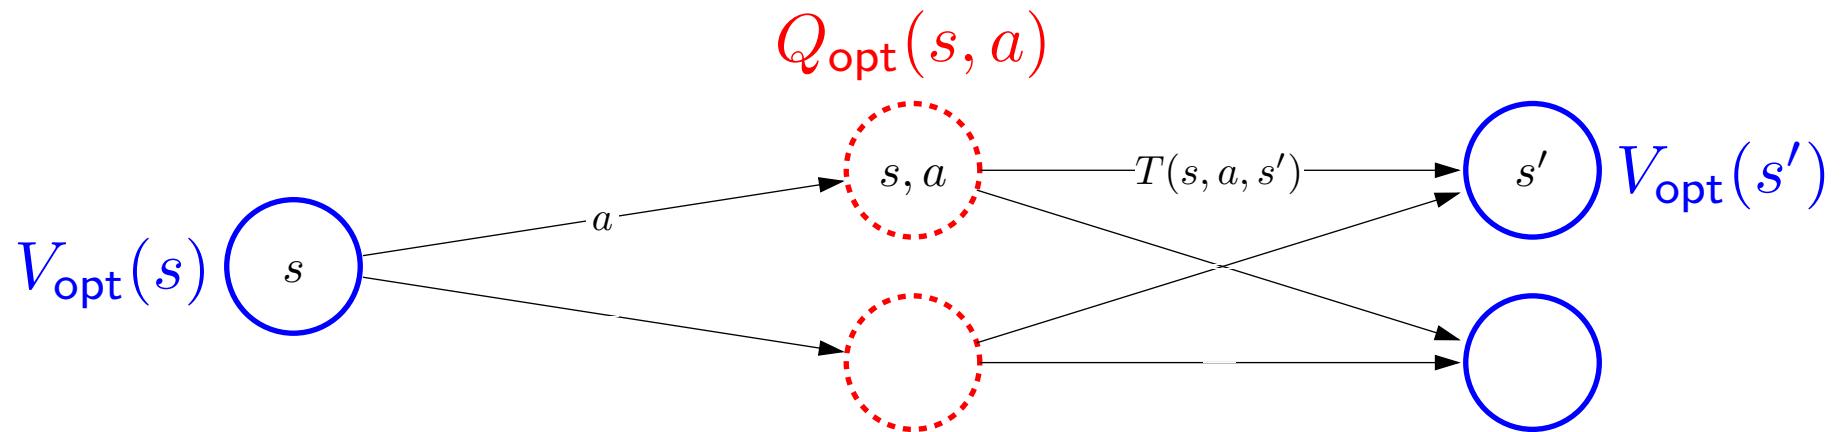
$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')].$$

Optimal value from state s :

$$V_{\text{opt}}(s) = \begin{cases} 0 & \text{if } \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} Q_{\text{opt}}(s, a) & \text{otherwise.} \end{cases}$$

- The recurrences for V_{opt} and Q_{opt} are identical to the ones for policy evaluation with one difference: in computing V_{opt} , instead of taking the action from the fixed policy π , we take the best action, the one that results in the largest $Q_{\text{opt}}(s, a)$.

Optimal policies



Note: value iteration recurrences mirror policy evaluation recurrences

Main difference: take **max** over actions a rather than using $a = \pi(s)$

- This looks kind of like sliding a bit of policy improvement into policy evaluation, and in fact that's exactly what it is. Value iteration can be interpreted as policy iteration where we only do one iteration of policy evaluation in the inner loop. The intuition is that there's no need to compute the value of a policy exactly when we're just going to change the policy anyway; we just need a rough estimate.

Value iteration



Algorithm: value iteration [Bellman, 1957]

Initialize $V_{\text{opt}}^{(0)}(s) \leftarrow 0$ for all states s .

For iteration $t = 1, \dots, t_{\text{VI}}$:

For each state s :

$$V_{\text{opt}}^{(t)}(s) \leftarrow \max_{a \in \text{Actions}(s)} \underbrace{\sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}^{(t-1)}(s')]}_{Q_{\text{opt}}^{(t-1)}(s, a)}$$

Time: $O(t_{\text{VI}}SAS')$

Interpretation: value iteration is equivalent to policy iteration where policy evaluation is run for one iteration

- By now, you should be able to go from recurrences to algorithms easily. Following the recipe, we simply iterate some number of iterations, go through each state s and then replace the equality in the recurrence with the assignment operator.
- Value iteration is also guaranteed to converge to the optimal value.
- What about the optimal policy? We get it as a byproduct. The optimal value $V_{\text{opt}}(s)$ is computed by taking a max over actions. If we take the argmax, then we get the optimal policy $\pi_{\text{opt}}(s)$.

Value iteration: dice game

s	end	in
$V_{\text{opt}}^{(t)}$	0.00	12.00 ($t = 100$ iterations)
$\pi_{\text{opt}}(s)$	-	stay

- Let us demonstrate value iteration on the dice game. Initially, the optimal policy is "quit", but as we run value iteration longer, it switches to "stay".

Value iteration: volcano crossing

Run

(or press ctrl-enter)

		-50	20
		-50	
2			

- As another example, consider the volcano crossing. Initially, the optimal policy and value correspond to going to the safe and boring 2. But as you increase numIters, notice how the value of the far away 20 propagates across the grid to the starting point.
- To see this propagation even more clearly, set slipProb to 0.

Convergence



Theorem: convergence

Suppose either

- discount $\gamma < 1$, or
- MDP graph is acyclic.

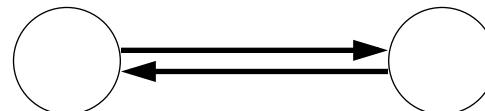
Then value iteration and policy iteration both converge to the correct answer.

Intuition: game will stop eventually



Example: non-convergence

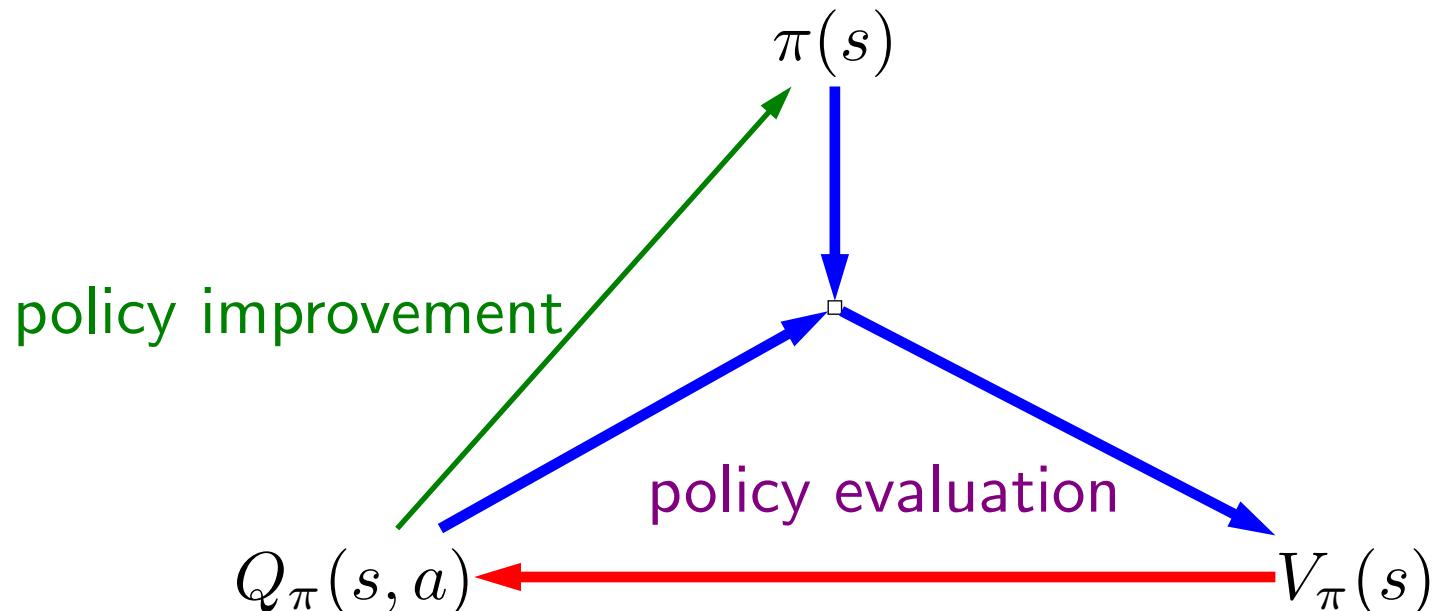
discount $\gamma = 1$, zero rewards



- Let us state more formally the conditions under which any of these algorithms that we talked about will work. A sufficient condition is that either the discount γ must be strictly less than 1 or the MDP graph is acyclic.
- We can reinterpret the discount $\gamma < 1$ condition as introducing a new transition from each state to a special end state with probability γ , multiplying all the other transition probabilities by $1 - \gamma$, and setting the discount $\gamma = 1$. The interpretation is that with probability γ , the game terminates at any state.
- In this view, we just need that a sampled path be finite with probability 1.
- We won't prove this theorem, but will instead give a counterexample to show that things can go badly if we have a cyclic graph and $\gamma = 1$. In the graph, whatever we initialize value iteration, value iteration will terminate immediately with the same value. In some sense, this isn't really the fault of value iteration, but it's because all paths are of infinite length. In some sense, if you were to play this game, you would never terminate, so we would never find out what your utility was at the end.



Connections



Methods

Policy evaluation $(\rightarrow \rightarrow)^*$

Policy improvement \rightarrow

Policy iteration $((\rightarrow \rightarrow)^* \rightarrow)^*$

Value iteration $(\rightarrow \rightarrow \rightarrow)^*$

Unifying idea

Algorithms:

- Search DP computes $\text{FutureCost}(s)$
- Policy evaluation computes policy value $V_\pi(s)$
- Value iteration computes optimal value $V_{\text{opt}}(s)$

Recipe:

- Write down recurrence: $f(s) = \dots f(s') \dots$
- Turn into iterative algorithm (replace mathematical equality with assignment operator)

- There are two key ideas in this lecture. First, the policy π , value V_π , and Q-value Q_π are the three key quantities of MDPs, and they are related via a number of recurrences which can be easily gotten by just thinking about their interpretations.
- Second, given recurrences that should hold on the thing that you're trying to compute, it's easy to turn these recurrences into algorithms that just iterate on those recurrences until convergence.



Summary

- **Markov decision processes** (MDPs) cope with uncertainty
- Solutions are **policies** rather than paths
- **Policy evaluation** computes policy value (expected utility)
- **Policy iteration** and **value iteration** computes optimal value (maximum expected utility) and optimal policy
- Main technique: write recurrences → algorithm
- Next time: reinforcement learning — when we don't know rewards, transition probabilities



cs221.stanford.edu/q

Question

What did you take away from this lecture?